



Developing Anthropomorphic Algorithms

~~FOR A LONG TIME, sociotechnical systems analysis within~~ computing science has developed formalisms of human-like traits, such as trust and comfort. These allow an intelligent agent to interact with other agents in its environment in ~~measured, cautious~~ ways: they might be used, for example, to decide whether it should accept information from another agent if its behaviour is becoming erratic (or to discard previous data which is no longer “trustworthy”).

ANTHROPOMORPHIC ALGORITHMS are algorithms which simulate a social human trait in an artificial agent. Examples of anthropomorphic algorithms already widely researched would be ones termed “computational trust formalisms”, versions of which are now decades old¹, but are rarely explored outside sociological and sociotechnical research.

These anthropomorphic algorithms are undergoing continual improvement², but some problems remain unexplored. Exploration of the techniques within the scope of AI development might yield solutions to problems of corrigibility³ and reward hacking⁴ in AI safety research. The technique implies that human behaviours can apply to non-biological “minds”, useful to philosophy of mind research⁵.

Within the realm of computing science research, more anthropomorphic algorithms are yet to be developed, helping fields such as human-computer interaction to create more authentically human-feeling interfaces. Anthropomorphic Algorithms are an important area of research which can be expected to flourish in both computer science and interdisciplinary research.

Future Developments

ONE PREREQUISITE TO THIS RESEARCH is a develop a software architecture which allows for the combination of multiple traits. Current methods for designing multiple traits would involve developing a single formalism which modelled multiple traits. However, an architecture which combined these traits would permit combining sociological and psychological theories which have already been formalised, without the additional work of creating an overarching formalism for every trait combination.

COMBINING SEVERAL TRAITS has a great deal of practical utility. To illustrate, one can imagine designing an interface to a mobile phone

¹ Stephen Paul Marsh. Formalising Trust as a Computational Concept. *Computing*, Doctor of(April):184, 1994. DOI: 10.2165/00128413-199409230-00010

² Seifeddine Kramdi. A modal approach to model computational trust. *PhD Thesis*, 2015. URL <https://tel.archives-ouvertes.fr/tel-01328169>; and Joana Urbano, Ana Paula Rocha, and Eugénio Oliveira. An approach to computational social trust. 27:113–131, 2014. DOI: 10.3233/AIC-130587

³ Stuart Armstrong, Nate Soares, Benja Fallenstein, and Eliezer Yudkowsky. Corrigibility. *AAAI Workshops: Workshops at the Twenty-Ninth AAAI Conference on Artificial Intelligence*, 2015

⁴ Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in AI safety. *CoRR*, abs/1606.06565, 2016. URL <http://arxiv.org/abs/1606.06565>

⁵ Aaron Sloman. The Structure of the Space of Possible Minds. pages 35–42, 1984

which takes into account a device's "feeling" of trust and comfort — two traits which have already been formalised into anthropomorphic algorithms.⁶ A mobile phone might have a degree of trust in the person it identifies as using it; less trustworthy users might be prohibited from accessing more sensitive information, such as medical information stored in systems like Apple's HealthKit database.

The device might also have a degree of comfort which is diminished when the user it identifies as using it begins acting erratically. One would expect the degree of trust it had in a previously trustworthy user to decrease as a result — perhaps its initial assessment of the user's trustworthiness was mistaken, or perhaps its identification of its user is incorrect. If erratic behaviour in this case decreased a feeling of trust, then trust and comfort are in some way linked.

One can imagine similar situations where trust might affect the device's degree of comfort, where switching from a trusted to an untrusted user might result in a sharp decrease in comfort. It is important to design effective and simple ways to model these and other situations, then, so as to make the engineering of these useful systems as simple as possible. Research — in a range of fields anthropomorphic algorithms touch, including and extending beyond computing science — would be affected too, as the construction and testing of anthropomorphic systems would be simplified if their creation can be simplified.

THE NEED for this system, which involves the integration of several traits, should be done via an agreed-upon software architecture and a formalised method for creating these formalisms. An anthropomorphic architecture would permit the combination of several systems, without the need to develop a new formalism for each combination.

A good architecture would also explore the possibilities for a general format of these formalisms. For a formalism to work in the architecture, it might need to adhere to certain requirements. This lends an opportunity to develop a suitable architecture for the formalisms themselves, providing guidelines for new formalisms to be constructed against and adding coherency to the growing inventory of formalisms being studied.

I propose that this work is a suitably sized and impacting topic for PhD level research. I also believe that the project holds much value, due to the pressing need for the anthropomorphic architecture to be developed. I am also inspired, however, due to the fact that my own masters level research involves developing new anthropomorphic algorithms for responsibility, and that the existence of this architecture would permit exciting new research opportunities is an observation born of my own enthusiasm for the topic.

⁶ Stephen Marsh, Pamela Briggs, Khalil El-Khatib, Babak Esfandiari, and John A. Stewart. Defining and Investigating Device Comfort. *Journal of Information Processing*, 19(7):231–252, 2011. ISSN 1882-6652. DOI: 10.2197/ip-sjip.19.231. URL <http://ci.nii.ac.jp/naid/110008508036/en/>; and Stephen Paul Marsh. Formalising Trust as a Computational Concept. *Computing, Doctor of (April)*:184, 1994. DOI: 10.2165/00128413-199409230-00010



Suitability

HAVING DEVELOPED the only existing anthropomorphic algorithm for responsibility formalism, as well as having experience in sociotechnical systems research, I am very well suited to pursue this particular research project. My honours year dissertation⁷ developed new sociotechnical modelling techniques for introducing human-like variance to workflow modelling with minimal overhead, so as to create human-readable workflow models which dynamically introduced human error during runtime. The models, and the dissertation project, were pure Python code. A paper on the work is currently in progress.⁸ Sociotechnical modelling tackles similar problems to anthropomorphic algorithms, and my familiarity with computational models of human traits shows my familiarity with the field.

FOLLOWING THIS, my MSci research project is to create a responsibility formalism — all references here to responsibility formalisms are my own work. I am excited to continue the work, and know that once properly developed as a research field, responsibility formalism will play a vibrant role in the creation of mixed-trait models. My familiarity with the state-of-the-art, and my proven ability to push the state of the art in anthropomorphic algorithms research, is a testament to my suitability in carrying out this important next development in the field.

MY CONFIDENCE IN MY FIELD is evidenced by my activity in the area. I have in-progress articles on the impact anthropomorphic algorithms can have in the philosophical arena, and will be talking on anthropomorphic algorithms' interdisciplinary potential at the "Let's Talk About [X]" undergraduate research conference in February 2017. My willingness to contribute to the field even in my spare time, and to speak publicly on the research as it stands, shows my dedication to research in anthropomorphic algorithms as a field. My confidence and experience in software engineering for sociotechnical systems, and my familiarity with anthropomorphic algorithms as a field, primes me to continue this research in a unique way.

Thank you for your consideration.

Is this section needed?

⁷ My honours year dissertation received an A2, and the award for the best software engineering of my year — despite it being a research project specifically instead of the engineering projects it competed with.

⁸ Tom Wallis and Tim Storer. Simulating variance in socio-technical behaviours using executable workflow fuzzing. Paper on the dynamic human-like variance of a workflow model, complete with a suitable modelling framework in Python, currently in progress. Recent draft can be found at <http://github.com/probablytom/fuzzimoss-paper>, 2016